

User Research at Fedi: Strategic Assessment & Recommendations

Independent Report — March 2026

The Short Version

Fedi recently completed a study to evaluate three onboarding flows for the wallet experience. The study produced a clear directional signal: the majority of participants preferred a simplified flow that gives users a meaningful choice before asking them to commit. That signal is credible and worth acting on.

But the more important story is not which flow won. It is what this study reveals about how Fedi currently approaches user research — and what is possible if that approach is upgraded.

What the Research Tells Us

The study asked 13 participants — a mix of internal team members, power users, and community leaders across LatAm, Africa, and Asia — to evaluate three onboarding designs and share their preferences.

The finding is consistent: most participants preferred a flow that presents a simple, clear choice upfront (auto-select or manual wallet setup) rather than either overwhelming them with options immediately or pre-selecting a provider before they have any context. The reasoning participants gave — across different regions and user types — was grounded in real dynamics: clarity, trust, and the feeling of having chosen rather than been pushed.

One piece of context worth noting: Option A in this study is not a novel concept — it is Fedi's current live onboarding flow. The majority of participants, particularly internal team members and community power users, had already encountered it before the study. The fact that Option B was preferred over something participants already knew and had used makes that preference a stronger signal than it might initially appear. Choosing the new over the familiar — without being primed to do so — suggests the preference reflects genuine usability improvement, not novelty.

These themes hold up under scrutiny. Terminology confusion (“wallet provider,” “federation,” “auto-select”) is a real and consistent barrier. Cognitive overload at the point of first decision is a real problem. The relationship between perceived consent and trust is a real UX dynamic — and one that matters especially for a product targeting communities that have historical reasons to be sceptical of financial systems that act on their behalf without asking.

The research points in a clear direction. It is a good starting point.

What the Research Cannot Tell Us

The study captures what people *say* they prefer. It does not capture what they *do*.

There is a well-documented gap between stated preference and actual behaviour in product research — and it is wider in fintech than almost anywhere else. A user in a moderated session with a researcher watching them navigate a prototype is not the same as a user in Jakarta or Nairobi downloading an app for the first time, alone, with patchy connectivity, and no one to ask for help.

The study also has structural limitations that a more experienced research practitioner would design out: all three flows were shown in the same fixed order across every session, which likely inflated preference for the option shown last. More than half the participants were internal team members, who understand federations, trust the organisation, and cannot fully simulate what it feels like to encounter these concepts for the first time. And because Option A is the current live app, the study was not comparing three equally unfamiliar concepts — it was comparing two new designs against something many participants already used daily, without acknowledging that asymmetry in how findings were interpreted. These are not criticisms of the people who ran the sessions — they are design problems in the study itself, and they are fixable.

The result is a signal. A strong enough signal to act on, but not strong enough to close the conversation. The right response is to move forward with the leading option while building in the kind of behavioural testing that confirms — or challenges — what the attitudinal study suggested.

The Bigger Pattern

What this study reflects is a wider dynamic that many product teams encounter: research gets done reactively, under time pressure, to resolve a decision that has already been partially made. The question being asked is “which of these three options should we build?” rather than “what does a user in this context actually need, and have we designed for it?”

That is not a criticism of intent. It is a description of what happens when research is treated as a deliverable rather than a discipline — when the goal is to produce an answer rather than to reduce uncertainty.

The engineering and product tension that has driven this study is a symptom of the same pattern. When teams disagree about the right direction and there is no systematic research infrastructure to generate shared evidence, opinions

fill the gap. Research becomes a tiebreaker rather than a foundation. It gets commissioned when a stalemate needs resolving, not before the decisions that create stalemates are made.

This is a solvable problem. But solving it requires treating research as a strategic function, not a project-by-project service.

The Analogy That Fits

In Bitcoin custody, there is a useful distinction between an **auditor** and a **custody architect**.

An auditor tells you the current state of a wallet — what's in it, whether the keys are intact, whether the last transaction checks out. It is a snapshot. Useful, necessary, but backward-looking.

A custody architect designs the system that surrounds the wallet: the signing policy, the key management structure, the recovery path, the human protocols that determine what happens when something goes wrong. Bitcoin is sound. But plenty of people have lost funds not because the protocol failed — it never does — but because nobody thought carefully enough about the human layer around it. The custody architecture is where experience, trust, and real-world behaviour matter most.

Commissioning one-off research studies when a decision needs making is auditing. It tells you the current state of user opinion and moves on. What Fedi needs — given where it is in its growth, the communities it is trying to serve, and the complexity of the concepts it is asking new users to navigate — is a custody architect for the user experience. Someone who designs the research infrastructure: what gets studied, when, how findings feed back into product decisions, and how that system evolves as the product does.

That is not a research monkey running studies on request. It is a research leader building the function.

What a Research-Led Approach Looks Like

The difference between reactive and embedded research is largely a matter of timing and structure.

Reactive research (current state): A design question needs an answer. A study is commissioned. Findings are delivered. The team acts on them or doesn't. The next decision starts the cycle again from scratch.

Embedded research (recommended state): Research is mapped to the product roadmap. Before a feature enters design, the relevant user questions have

already been identified and, where possible, answered. Development teams are building on evidence rather than assumption. Post-launch, behavioural data feeds back into the next research cycle.

Concretely, this means building a **UX research roadmap** aligned to Fedi's product development calendar — identifying the 6–12 months of planned builds, mapping the user knowledge gaps associated with each, and scheduling the right type of research (generative, evaluative, or behavioural) at the right point in the build cycle.

For a product serving communities in emerging markets where digital financial literacy varies enormously, where multiple languages are in play, and where trust is hard-won and easily broken, this kind of intentional research infrastructure is not a luxury. It is how you build something that actually works for the people it is meant to serve — and how you know when to change course before a full build has been committed to.

Recommended Next Steps

Immediately: - Treat the current study's findings as a working hypothesis, not a verdict. Proceed with the preferred onboarding direction, but build in unmoderated task testing and behavioural metrics as the flow is developed. - Establish a feedback loop between research and engineering — so that when behavioural data comes in, there is a clear process for incorporating it rather than re-opening design debates.

Within 90 days: - Map the product roadmap against a set of open user questions — what do we not yet know about users that could affect the success of each planned build? - Prioritise the top three research questions and design studies that address them *before* the associated builds begin. - Increase the proportion of genuine new users — particularly in LatAm, Africa, and South-east Asia — in any future research. Internal participants and power users have a role, but not as primary evidence.

Structurally: - Define what a research function at Fedi looks like: not a single study, but a practice. Who owns it, how it connects to product decisions, and how findings are documented and built upon over time. - Invest in that function with the same seriousness applied to engineering infrastructure. The human layer of a financial product is as consequential as the technical layer — and just as dependent on systematic, rigorous attention.

Closing Thought

The onboarding study is a good start. It asked the right questions, surfaced real themes, and produced a direction worth pursuing. The goal of this report is not to diminish that work — it is to show what becomes possible when it is part of a larger, more intentional research practice rather than a one-time exercise.

Fedi is building financial infrastructure for communities that have rarely been well-served by it. Getting the human experience right is not separate from that mission. It is the mission.

Independent report prepared March 2026. Based on review of session transcripts, research notes, and study design documentation provided.